



directly, and then used its knowledge of page numbers to inspect page 5 accurately.

Now for the final experiment of Day 4: testing what happens with scanned (image-only) PDFs.

Milestone 7: The Scanned PDF Experiment

The "Why"

As we discussed in the mental model step, PyMuPDF performs **text extraction** (pulling embedded text characters). It does **not** perform OCR (Optical Character Recognition).

To solidify this distinction before writing your journal entry, we want to observe what `read_pdf` returns when given an image-based PDF.

Your Instructions

1. **Option A (Quickest):** Take any image file on your computer (`.png` or `.jpg`), or take a screenshot, open it, and print/save it as a PDF named `scanned_test.pdf`.
2. **Option B:** Find any scanned document PDF online.
3. Move `scanned_test.pdf` into your `tool-use-agent-week4` folder.
4. Run `python pdf_chat.py` and give it `scanned_test.pdf`.
5. Observe the warning message that our tool outputs when no text characters are found!

Your Turn

1. Run `pdf_chat.py` against a scanned/image PDF.
2. Reply with "Done" and briefly share what output or warning message was printed when you tried to read it!

👤 🗨️ ↻ 📄 ...

+ Ask Gemini

Flash ▾

